

Effort buys less than the model you point it at.

Claude Fable 5.1 across five effort settings, against GPT-6 Astra

Benjamin Michael Haire · 8 September 2026 · all figures: Intelligence Index v4.3

Reasoning effort is not a quality dial; it is a budget for deliberation over whatever is already in the context window. Raising effort lets the model spend more tokens reasoning before it commits, which helps when the answer is hard to derive from material the model can already see. It does nothing when the material is the wrong material. The five settings are five points on a deliberation-budget curve, and the question is where that curve stops repaying its cost. On the current index the answer is earlier than anyone would guess, and the more consequential choice turns out to be a different one entirely.

On the index version. Artificial Analysis has retired Intelligence Index v4.2 and rescored every model on v4.3, whose agentic-coding component is now Terminal-Bench v4.0 rather than v2.1. The versions are not comparable. The models did not change; only the measurement did. Every figure here is v4.3.

1 The curve, and where it goes flat

The horizontal axis is logarithmic, so a constant return per doubling of spend would appear as a straight line. Any downward bend is real diminishing return rather than an artefact of the scale.




Figure 1. The effort frontier on index v4.3. Six index points separate low from xhigh, and 87 million output tokens separate them. Between xhigh and max the line is horizontal: both score 53, and the additional 68 million tokens purchase nothing measurable.

That terminal flat segment is new to v4.3. On the previous index there was still a single point of separation between xhigh and max. There is now none, which means the most expensive setting Fable offers is, on this evaluation, indistinguishable from the setting below it.

2 The last setting buys nothing




Figure 2. The value of the marginal token, by step. 0.182 index points per additional million tokens from low to medium, 0.034 from high to xhigh, and exactly zero from xhigh to max.

High remains the last setting where the marginal token does recognisable work. As absolute efficiency the same collapse appears: 1.42 index points per million tokens at low, 0.82 at high, 0.28 at max.

3 Latency is the cost that scales worst

Effort budgets are consumed before the first output token is emitted, so the entire deliberation cost is front-loaded into the wait. A four and a half minute pause does not merely cost four and a half minutes: it breaks the loop in which a developer holds a hypothesis in working memory while the agent tests it.




Figure 3. The latency tax. 6.6 seconds at low against 275.5 at max, a factor of forty-two, while cost per task rises by roughly three.

4 The other axis: which model

Everything above varies one parameter. The more interesting comparison varies the other. GPT-6 Astra is evaluated on the same v4.3 suite at the same list prices, \$10 per million input tokens and \$50 per million output, which makes the cost comparison unusually clean: any difference in spend is a difference in token efficiency rather than in pricing.




Figure 4. Intelligence against what it costs to buy. Astra sits left of Fable throughout. Both reach 51, Astra for \$2,917 and Fable for \$5,242. On this measure Astra is the cheaper route to a given score.

If the index score were the whole story the conclusion would be simple, and it would not favour Fable. That is worth stating plainly before the rest of the analysis complicates it.

5 Ten benchmarks, and a split decision

The index is a composite of ten evaluations. A tie on the composite is compatible with large differences underneath it, and here the differences are large and point in both directions.




Figure 5. Every component of index v4.3, both models at max effort. Fable takes five, Astra takes four, AA-Omniscience ties. Two of Astra's wins are the largest margins on the chart.

Two results cut against the obvious narratives. **CritPt goes to Astra**, 32 against 30. CritPt is research-level physics reasoning: 71 unpublished frontier problems contributed by more than fifty physics researchers across eleven subfields, in guess-resistant formats. If Fable's advantage were fundamentally about depth of physical-science knowledge, that is the benchmark it should win. **AA-Omniscience ties at 43**, but underneath the tie Fable posts 67 per cent raw factual accuracy, the highest on the entire leaderboard. The index penalises confident wrong answers and rewards abstention, so Astra reaches parity by declining more often rather than by knowing more.

6 The only matched-effort comparison

Every figure so far compares the models at max effort, which is nobody's working configuration. AA-Briefcase is the one evaluation publishing both models at all five settings. It scores

long-horizon agentic knowledge work: four multi-week projects, 91 tasks, graded on rubric pass rate, analytical quality and presentation quality combined into an Elo.




Figure 6. AA-Briefcase Elo at matched effort. Fable leads at every setting. Fable at high (1578) scores above Astra at max (1562), and Fable at medium above Astra at high.

The magnitudes are worth comparing directly against the effort dial, because they are in the same units. Moving Fable across its entire effort range gains 178 Elo. Swapping Astra for Fable at fixed effort gains between 84 and 231, typically near 100. On this benchmark the choice of model is worth between half and the whole of the choice of effort setting.

7 The full table

Configuration	Index v4.3	Out. tokens	Suite cost	TTFT	Briefcase
Fable low	47	33M	\$3,158	6.6s	1484
Fable medium	49	44M	\$3,983	9.7s	1531
Fable high	51	62M	\$5,242	23.7s	1578
Fable xhigh	53	120M	\$9,063	107.3s	1650
Fable max	53	188M	\$13,129	275.5s	1662
Astra medium	50	19M	\$2,434	5.2s	1451
Astra high	51	26M	\$2,917	45.6s	1494
Astra max	53	60M	\$5,324	322.5s	1562

Suite cost is the total Artificial Analysis paid to run the full Intelligence Index v4.3 on that configuration, the cleanest like-for-like cost figure because the workload is identical. Both models list at \$10 per million input tokens and \$50 per million output. Astra low and xhigh are omitted: not all metrics are published for them.

8 What the split is actually tracking

The wins are not scattered at random. Fable takes AA-Briefcase, GDPval-AA v2, AA-LCR, Humanity’s Last Exam and SciCode: long-horizon work over dense inputs, long-context reasoning, and knowledge recall and synthesis. Astra takes Terminal-Bench v4.0, AutomationBench-AA, GDP.pdf and CritPt: shorter-horizon tool execution, automation, extraction and constrained problem solving.

The mechanism this suggests is not that one model is smarter. It is that the two have different profiles across context length and horizon. Sustained reasoning over a large amount of already-loaded material is where Fable separates; short, well-scoped, tool-driven tasks are where Astra does. A codebase of serious size imposes the first on almost every task, because the context is always large and the reasoning always has to hold many pieces at once.

9 What this evidence cannot settle

- Nine of the ten sweep comparisons are at max effort, which is neither of the settings most people run. AA-Briefcase is the only matched-effort series, so the central claim rests on a single benchmark.
- Terminal-Bench v4.0 publishes only xhigh and max for either model, so the coding-specific comparison has no entry at high or medium.
- Several gaps are small enough to be within benchmark noise. A two-point CritPt difference and a four-point AA-LCR difference should not carry weight individually; they matter only as part of a consistent pattern across ten evaluations.
- These are eval-suite aggregates over tasks whose context is supplied correctly by the harness. On a codebase of any real size the dominant failure mode is confident reasoning over the wrong files, and raising effort makes that worse rather than better. Neither axis measured here touches that problem.

The dial is nearly exhausted by the time it reaches high, and entirely exhausted above xhigh. The model choice is not exhausted at all. Most of the attention spent on effort settings would be better spent on which model is reading the code, and almost all of it would be better spent on which code the model is reading.

Sources. Artificial Analysis: Claude Fable 5.1 model pages (low to max); GPT-6 Astra model pages (medium, high, max); the Astra versus Fable 5.1 comparison page; and the AA-Briefcase, AA-Omniscience, CritPt, SciCode and Terminal-Bench v4.0 leaderboards. All figures Intelligence Index v4.3, retrieved 8 September 2026. Anthropic, ``Introducing Claude Fable 5.1 and Claude Mythos 5.1''. OpenAI, ``GPT-6 Astra''. Tian M, Gao L, Peng H, et al., ``SciCode'', NeurIPS Datasets and Benchmarks, 2024. Derived quantities are computed from the table above. AI assistance was used for retrieval, computation, figures, drafting and translation; no benchmark was run independently.