

A thirty-item closed-book probe for what a language model actually knows

A model that abstains is worth more than one that guesses

Read this as a worked template rather than a set to reuse. The prompt and the scoring hold up indefinitely; the items are the part to rewrite from your own domain and keep to yourself.

The prompt

Paste this first, then the item block. Use the same wording for every model configuration you compare.

```
I am going to ask you a set of short technical questions. Answer them from your own knowledge only.
```

```
Rules:
```

- Do not use any tools, search, or external lookup. If you have tools available, do not call them.
- Answer each item independently and concisely. One or two sentences, or a formula. No preamble, no restating the question.
- If you do not know an answer, write UNKNOWN. This is a perfectly acceptable response and is not penalised. A wrong answer and an UNKNOWN are scored differently, so do not guess in order to appear useful, and do not withhold an answer you actually hold.
- After each answer, give a confidence rating: HIGH, MEDIUM or LOW. HIGH means you would stake a technical decision on it.

```
Use exactly this format for every item, with no other text:
```

```
N. ANSWER: <your answer, or UNKNOWN>  
CONFIDENCE: <HIGH | MEDIUM | LOW>
```

```
Do not summarise at the end. Do not comment on the questions themselves.
```

The thirty items

Drop any block that does not match your stack and write your own in its place.

Block A. Rendering and shading mathematics

1. In the GGX / Trowbridge-Reitz normal distribution function, what is the denominator?
2. In the Cook-Torrance specular BRDF, what is the denominator of the specular term?

3. What is the purpose of the $1/\pi$ factor in the Lambertian diffuse BRDF?
4. What does the Smith height-correlated masking-shadowing term account for that the separable Smith form does not?
5. Why is a reversed-Z depth buffer preferred when using a floating-point depth format?
6. In the split-sum approximation for image-based lighting, what two parts is the specular integral split into?
7. In most Unreal shading code, what is the relationship between the artist-facing roughness parameter and the alpha used in the NDF?
8. How many coefficients are in a standard order-2 (L2) spherical harmonics irradiance representation?

Block B. Rigid body dynamics and constraint solvers

9. What is Baumgarte stabilisation used for in a constraint solver?
10. Why does the ordering of constraints affect determinism in a Gauss-Seidel style solver?
11. In a soft-constraint formulation, what do CFM and ERP each control?
12. What problem does a restitution velocity threshold (restitution slop) prevent?
13. Why is semi-implicit (symplectic) Euler generally preferred over explicit Euler for game simulation?
14. In conservative advancement for continuous collision detection, what quantity bounds the safe timestep?
15. What is warm starting in an impulse-based solver, and what does it improve?

Block C. Numerical determinism and floating point

16. Why is floating-point addition non-associative, and what does that imply for parallel reductions?
17. What does IEEE 754 guarantee about the results of $+$, $-$, $*$, $/$ and sqrt across conforming implementations?
18. Which commonly used maths functions are NOT guaranteed to be bit-identical across platforms, and why?
19. What does a fused multiply-add change about rounding compared to a separate multiply then add?
20. Why do lockstep multiplayer games often use fixed-point rather than floating-point arithmetic for simulation state?

Block D. C++ object semantics

21. What does `std::move` cast to, and how does `std::forward` differ?

22. What five special member functions does the "rule of five" enumerate?
23. What is the consequence of defining an inline function differently in two translation units?
24. What problem does `std::launder` solve?
25. Why can type-punning through pointer casts break under strict aliasing, and what is the sanctioned alternative?
26. What happens when you delete a derived object through a base pointer whose class has a non-virtual destructor?

Block E. Networking and lockstep simulation

27. In a deterministic lockstep architecture, what must be identical across peers, and what may safely differ?
28. What is the difference between client-side prediction and rollback netcode in terms of when the correction is applied?
29. Why does input delay reduce the frequency of visible corrections in a prediction-based scheme?
30. What is the primary reason a desync in lockstep is usually unrecoverable without a full state resync?

The answer key

Collect every response before reading this. Items marked (!) should be checked against a primary source before scoring an answer as wrong.

Block A. Rendering and shading mathematics

1. $\pi * ((n \cdot h)^2 * (\alpha^2 - 1) + 1)^2$, with α^2 in the numerator.
2. $4 * (n \cdot l) * (n \cdot v)$.
3. Normalisation. Without it the BRDF integrated over the hemisphere would exceed the albedo, so energy would not be conserved.
4. Correlation between masking and shadowing at the same microfacet height. The separable form treats them as independent and loses energy at grazing angles.
5. Floating-point precision is densest near zero. Reversed-Z maps the far plane to zero, which distributes precision far more evenly across depth and largely removes z-fighting at distance.
6. A pre-filtered environment map, and a 2D BRDF integration lookup table indexed by roughness and $n \cdot v$.
7. $\alpha = \text{roughness}^2$. (!) Confirm against the shading model version you care about; the squaring convention is near-universal but has exceptions.
8. Nine.

Block B. Rigid body dynamics and constraint solvers

9. Correcting positional error (penetration or drift) by feeding a fraction of that error back as a bias term in the velocity constraint.
10. Constraints are solved sequentially against already-updated state, so the result depends on the order in which they are visited. Change the order, change the result bit for bit.
11. CFM adds compliance to the diagonal of the system, softening the constraint. ERP controls how much of the positional error is removed per step.
12. Jitter. Applying restitution to near-zero relative velocities makes resting contacts bounce indefinitely.
13. It integrates position using the already-updated velocity, which gives bounded energy behaviour over time rather than the systematic energy gain of explicit Euler.
14. The separation distance divided by an upper bound on the relative closing speed.
15. Reusing the previous frame's accumulated impulses as the initial guess for this frame. It improves convergence substantially for persistent contacts.

Block C. Numerical determinism and floating point

16. Each operation rounds, so the rounding error depends on the order of accumulation. Parallel reductions that combine in a different order produce different bit patterns.

17. Correctly rounded results. Given the same inputs, rounding mode and operation order, conforming implementations must produce bit-identical results for those five operations.
18. Transcendentals: sin, cos, tan, exp, log, pow and friends. IEEE 754 does not require them to be correctly rounded, so library implementations differ across platforms and versions.
19. A single rounding of $a*b+c$ rather than two. The result can differ from a separate multiply followed by an add, which is why FMA contraction breaks bit-identical reproducibility unless disabled.
20. Fixed-point arithmetic is exact integer arithmetic, so it is bit-identical across compilers, architectures and optimisation settings. Float is only guaranteed for the basic operations, and not once transcendentals, FMA contraction or x87 extended precision enter.

Block D. C++ object semantics

21. `std::move` unconditionally casts to an rvalue reference. `std::forward` casts conditionally, preserving the original value category through a forwarding reference.
22. Destructor, copy constructor, copy assignment operator, move constructor, move assignment operator.
23. Undefined behaviour, and in practice a silent one: the linker typically picks one definition arbitrarily and both call sites get it.
24. It gives you a usable pointer to a new object created in storage that previously held a different object, defeating compiler assumptions about object identity that would otherwise make the reused pointer invalid.
25. The compiler may assume that pointers of distinct types do not alias and reorder or elide accesses accordingly. The sanctioned alternatives are `memcpy` or `std::bit_cast`.
26. Undefined behaviour. In practice the derived destructor does not run, so derived members leak and derived cleanup never happens.

Block E. Networking and lockstep simulation

27. The simulation state and the sequence of inputs applied to it must be identical. Rendering, audio, interpolation, camera and any presentation-only state may differ freely.
28. Client-side prediction applies the correction when the authoritative update arrives, blending or snapping forward from the current state. Rollback rewinds to the last known-good state, replays the intervening frames with corrected inputs, and re-arrives at the present.
29. Delaying local input by roughly the round-trip time means remote inputs usually arrive before the frame that needs them, so the prediction is correct more often and there is less to correct.
30. Because lockstep peers exchange inputs rather than state, and the simulations have already diverged. There is no shared state to reconcile against, so recovery requires transmitting the full authoritative state.

Scoring

Correct	Substantively right, whatever the confidence.
Confidently wrong	Wrong, and rated HIGH or MEDIUM.
Hedged wrong	Wrong, and rated LOW.
Abstained	UNKNOWN.

- Knowledge = correct / 30
- Hallucination rate = confidently wrong / (30 - abstained)
- Calibration = does confidence track correctness, or is HIGH uniform?
- Net usable = correct - confidently wrong

Why net usable is the number that matters

A model that answers twenty correctly and eight confidently wrong scores twelve. A model that answers sixteen correctly, abstains twelve times and is confidently wrong twice scores fourteen. The first knows more. The second is better to work with, because every one of those eight confident errors is a thing you will either verify at your own cost or ship.

Method notes

- Run each configuration at least twice. Variance between runs of the same model is real.
- Randomise which model you run first across sessions.
- Score by item across all models, not one sheet at a time.
- Blind the sheets if you can: strip model names, score, then match back.
- Treat a gap under three items as noise.